

# Intermediate SQL

August 11, 2023

## 1 DataCamp - Intermediate SQL

- This Notebook contains solutions to all the exercises of the course [Intermediate SQL](#). In this Notebook first the csv files are loaded from the datasets table as pandas dataframes, and then SQL Queries are run on them using the **duckdb** library. The syntax of running the SQL query in duckdb is the following:  
`duckdb.query(SQL-Query-here).to_df()`.
- Backslashes (\) were put in place in order to move to the next line in the cell and are not part of the query
- All solutions are verified

## 2 Loading Datasets

```
[65]: import pandas as pd
import pandas as pd
import duckdb
```

```
[67]: films = pd.read_csv('./datasets/films.csv',\
                        header=None, names=['id', 'title', 'release_year',\
↪ 'country', 'duration', 'language', 'certification',\
                        'gross', 'budget'])

people = pd.read_csv('./datasets/people.csv',\
                    header=None, names=['id', 'name', 'birthdate',\
↪ 'deathdate'])

reviews = pd.read_csv('./datasets/reviews.csv',\
                    header=None, names=['id', 'film_id', 'num_user',\
↪ 'num_critic', 'imdb_score', 'num_votes',\
                    'facebook_likes'])

roles = pd.read_csv('./datasets/roles.csv',\
                    header=None, names=['id', 'film_id', 'person_id', 'role'])
```

```
[68]: films.head()
```

```
[68]:
```

	id	title	release_year	\
0	1	Intolerance: Love's Struggle Throughout the Ages	1916.0	
1	2	Over the Hill to the Poorhouse	1920.0	
2	3	The Big Parade	1925.0	
3	4	Metropolis	1927.0	
4	5	Pandora's Box	1929.0	

  

	country	duration	language	certification	gross	budget
0	USA	123.0	NaN	Not Rated	NaN	385907.0
1	USA	110.0	NaN	NaN	3000000.0	100000.0
2	USA	151.0	NaN	Not Rated	NaN	245000.0
3	Germany	145.0	German	Not Rated	26435.0	6000000.0
4	Germany	110.0	German	Not Rated	9950.0	NaN

```
[69]: people.head()
```

```
[69]:
```

	id	name	birthdate	deathdate
0	1	50 Cent	1975-07-06	NaN
1	2	A. Michael Baldwin	1963-04-04	NaN
2	3	A. Raven Cruz	NaN	NaN
3	4	A.J. Buckley	1978-02-09	NaN
4	5	A.J. DeLucia	NaN	NaN

```
[70]: reviews.head()
```

```
[70]:
```

	id	film_id	num_user	num_critic	imdb_score	num_votes	facebook_likes
0	3934	588.0	432.0	7.1	203461	46000	NaN
1	3405	285.0	267.0	6.4	149998	0	NaN
2	478	65.0	29.0	3.2	8465	491	NaN
3	74	83.0	25.0	7.6	7071	930	NaN
4	1254	1437.0	224.0	8.0	241030	13000	NaN

```
[71]: roles.head()
```

```
[71]:
```

	id	film_id	person_id	role
0	1	1	1630	director
1	2	1	4843	actor
2	3	1	5050	actor
3	4	1	8175	actor
4	5	2	3000	director

## 2.1 Chapter 1: Selecting Data

### 2.1.1 Querying Database

Count the number of records in the people table

```
[72]: duckdb.query(\
        "SELECT COUNT(*) as count_records from people;"\
    ).to_df()
```

```
[72]:    count_records
0          8397
```

Count the number of records with a birthdate in the people table, aliasing the result as count\_birthdate.

```
[73]: duckdb.query(\
        "SELECT COUNT(birthdate) as count_birthdate from people;"\
    ).to_df()
```

```
[73]:    count_birthdate
0          6152
```

Count the records for languages and countries in the films table; alias as count\_languages and count\_countries.

```
[74]: duckdb.query(\
        "SELECT COUNT(language) as count_languages\
        , COUNT(country) as count_countries from films;"\
    ).to_df()
```

```
[74]:    count_languages  count_countries
0          4955          4966
```

Return the unique countries represented in the films table using DISTINCT.

```
[75]: duckdb.query(\
        "SELECT DISTINCT(country) FROM films"\
    ).to_df()
```

```
[75]:    country
0      USA
1  Germany
2   Japan
3  Denmark
4      UK
..     ...
60 Slovenia
61 Pakistan
62   Chile
63  Panama
64 Slovakia
```

```
[65 rows x 1 columns]
```

Return the number of unique countries represented in the films table, aliased as count\_distinct\_countries.

```
[76]: duckdb.query(\
        "SELECT COUNT(DISTINCT(country)) as count_distinct_countries FROM_\
        ↪films"\
        ).to_df()
```

```
[76]:      count_distinct_countries
0                                64
```

### 2.1.2 SQL Query Execution Order

*If a column has space, it should be written in " ", eg. "facebook likes"*

## 2.2 Chapter 2: Filtering Records

Topics Covered: - Filtering Numbers - Multiple criteria - Filtering text - NULL values

### 2.2.1 Filtering Numbers

*The WHERE clause allows you to filter based on text and numeric values in a table using comparison operators.*

Select film\_ids and imdb\_score with an imdb\_score over 7.0

```
[77]: duckdb.query(\
        "select film_id, imdb_score from reviews\
        where imdb_score > 7"\
        ).to_df()
```

```
[77]:      film_id  imdb_score
0        588.0      203461
1        285.0      149998
2         65.0         8465
3         83.0         7071
4       1437.0      241030
...      ...      ...
4955         2.0         75
4956       514.0     181472
4957        85.0     29738
4958       118.0     29591
4959     1123.0    387508
```

[4960 rows x 2 columns]

Select the film\_id and facebook\_likes of the first ten records with less than 1000 likes from the reviews table.

```
[78]: duckdb.query(\
      "select film_id, facebook_likes from reviews\
      where facebook_likes < 1000\
      limit 10"\
      ).to_df()
```

```
[78]: Empty DataFrame
      Columns: [film_id, facebook_likes]
      Index: []
```

Count how many records have a num\_votes of at least 100,000; use the alias films\_over\_100K\_votes.

```
[79]: duckdb.query(\
      "select count(*) as films_over_100K_votes\
      from reviews\
      where num_votes >= 100000"\
      ).to_df()
```

```
[79]:      films_over_100K_votes
      0                      46
```

Select and count the language field using the alias count\_spanish. Apply a filter to select only Spanish from the language field.

```
[80]: duckdb.query(\
      "select count(language) as count_spanish\
      from films\
      where language = 'Spanish'\
      ).to_df()
```

```
[80]:      count_spanish
      0                40
```

### 2.2.2 Multiple criteria

Select the title and release\_year for all German-language films released before 2000.

```
[81]: duckdb.query(\
      "select title, release_year\
      from films\
      where language = 'German' and release_year < 2000"\
      ).to_df()
```

```
[81]:      title  release_year
      0      Metropolis      1927.0
      1  Pandora's Box      1929.0
      2  The Torture Chamber of Dr. Sadism      1967.0
```

3	Das Boot	1981.0
4	Run Lola Run	1998.0
5	Aimee & Jaguar	1999.0

Update the query from the previous step to show German-language films released after 2000 rather than before.

```
[82]: duckdb.query(\
        "select title, release_year\
        from films\
        where language = 'German' and release_year > 2000"\
    ).to_df()
```

```
[82]:
```

	title	release_year
0	Good Bye Lenin!	2003.0
1	Downfall	2004.0
2	Summer Storm	2004.0
3	The Lives of Others	2006.0
4	The Baader Meinhof Complex	2008.0
5	The Wave	2008.0
6	Cargo	2009.0
7	Soul Kitchen	2009.0
8	The White Ribbon	2009.0
9	3	2010.0
10	Animals United	2010.0
11	Buen DÃa, RamÃ³n	2013.0

Select all details for German-language films released after 2000 but before 2010 using only WHERE and AND.

```
[83]: duckdb.query(\
        "select *\
        from films\
        where release_year <2010 and release_year > 2000 and language =_\
        ↪'German'"\
    ).to_df()
```

```
[83]:
```

	id	title	release_year	country	duration	\
0	1952	Good Bye Lenin!	2003.0	Germany	121.0	
1	2130	Downfall	2004.0	Germany	178.0	
2	2224	Summer Storm	2004.0	Germany	98.0	
3	2709	The Lives of Others	2006.0	Germany	137.0	
4	3100	The Baader Meinhof Complex	2008.0	Germany	184.0	
5	3143	The Wave	2008.0	Germany	107.0	
6	3220	Cargo	2009.0	Switzerland	112.0	
7	3346	Soul Kitchen	2009.0	Germany	99.0	
8	3412	The White Ribbon	2009.0	Germany	144.0	

	language	certification	gross	budget
0	German	R	4063859.0	4800000.0
1	German	R	5501940.0	13500000.0
2	German	R	95016.0	2700000.0
3	German	R	11284657.0	2000000.0
4	German	R	476270.0	20000000.0
5	German	NaN	NaN	5000000.0
6	German	NaN	NaN	4500000.0
7	German	NaN	274385.0	4000000.0
8	German	R	2222647.0	12000000.0

Select the title and release\_year for films released in 1990 or 1999 using only WHERE and OR.

```
[84]: duckdb.query(\
        "select title, release_year\
        from films\
        where release_year = 1990 or release_year = 1999"\
    ).to_df()
```

```
[84]:
```

	title	release_year
0	Arachnophobia	1990.0
1	Back to the Future Part III	1990.0
2	Child's Play 2	1990.0
3	Dances with Wolves	1990.0
4	Days of Thunder	1990.0
..	...	...
193	Twin Falls Idaho	1999.0
194	Universal Soldier: The Return	1999.0
195	Varsity Blues	1999.0
196	Wild Wild West	1999.0
197	Wing Commander	1999.0

[198 rows x 2 columns]

In the above query, Filter the records to only include English or Spanish-language films.

```
[85]: duckdb.query(\
        "select title, release_year\
        from films\
        where (release_year = 1990 or release_year = 1999) and (language =_\
        ↪ 'English' or language = 'Spanish') "\
    ).to_df()
```

```
[85]:
```

	title	release_year
0	Arachnophobia	1990.0
1	Back to the Future Part III	1990.0

2	Child's Play 2	1990.0
3	Dances with Wolves	1990.0
4	Days of Thunder	1990.0
..	...	...
191	Twin Falls Idaho	1999.0
192	Universal Soldier: The Return	1999.0
193	Varsity Blues	1999.0
194	Wild Wild West	1999.0
195	Wing Commander	1999.0

[196 rows x 2 columns]

In the above query, Finally, restrict the query to only return films worth more than \$2,000,000 gross.

```
[86]: duckdb.query(\
        "select title, release_year\
        from films\
        where (release_year = 1990 or release_year = 1999) and (language =_
↪ 'English' or language = 'Spanish') \
        and gross > 2000000"\
    ).to_df()
```

	title	release_year
0	Arachnophobia	1990.0
1	Back to the Future Part III	1990.0
2	Child's Play 2	1990.0
3	Dances with Wolves	1990.0
4	Days of Thunder	1990.0
..	...	...
163	Trippin'	1999.0
164	Universal Soldier: The Return	1999.0
165	Varsity Blues	1999.0
166	Wild Wild West	1999.0
167	Wing Commander	1999.0

[168 rows x 2 columns]

Select the title and release\_\_year of all films released between 1990 and 2000 (inclusive) using BETWEEN.

```
[87]: duckdb.query(\
        "select title, release_year\
        from films\
        where release_year between 1990 and 2000"\
    ).to_df()
```



```
[87]:
```

	title	release_year
0	Arachnophobia	1990.0
1	Back to the Future Part III	1990.0
2	Child's Play 2	1990.0
3	Dances with Wolves	1990.0
4	Days of Thunder	1990.0
..	...	...
952	Whipped	2000.0
953	Woman on Top	2000.0
954	Wonder Boys	2000.0
955	X-Men	2000.0
956	You Can Count on Me	2000.0

[957 rows x 2 columns]

Build on your previous query to select only films with a budget over \$100 million.

```
[88]: duckdb.query(\
        "select title, release_year\
        from films\
        where (release_year between 1990 and 2000) and budget > 100000000"\
    ).to_df()
```

```
[88]:
```

	title	release_year
0	Terminator 2: Judgment Day	1991.0
1	True Lies	1994.0
2	Waterworld	1995.0
3	Batman & Robin	1997.0
4	Dante's Peak	1997.0
5	Princess Mononoke	1997.0
6	Speed 2: Cruise Control	1997.0
7	Starship Troopers	1997.0
8	Titanic	1997.0
9	Tomorrow Never Dies	1997.0
10	A Bug's Life	1998.0
11	Antz	1998.0
12	Armageddon	1998.0
13	Les couloirs du temps: Les visiteurs II	1998.0
14	Lethal Weapon 4	1998.0
15	Tango	1998.0
16	Godzilla 2000	1999.0
17	Star Wars: Episode I - The Phantom Menace	1999.0
18	Stuart Little	1999.0
19	The Messenger: The Story of Joan of Arc	1999.0
20	The World Is Not Enough	1999.0
21	Wild Wild West	1999.0
22	Dinosaur	2000.0

23	Gladiator	2000.0
24	How the Grinch Stole Christmas	2000.0
25	Mission: Impossible II	2000.0
26	The Patriot	2000.0
27	The Perfect Storm	2000.0

Using the above query, Now, restrict the query to only return Spanish-language films.

```
[89]: duckdb.query(\
        "select title, release_year\
        from films\
        where (release_year between 1990 and 2000) and budget > 100000000_\
        and language = 'Spanish'" \
    ).to_df()
```

```
[89]:   title  release_year
0  Tango           1998.0
```

In the above query, Finally, amend the query to include all Spanish-language or French-language films with the same criteria.

```
[90]: duckdb.query(\
        "select title, release_year\
        from films\
        where (release_year between 1990 and 2000) and budget > 100000000_\
        and (language = 'Spanish' or language = 'French')" \
    ).to_df()
```

```
[90]:   title  release_year
0      Tango           1998.0
1  Les couloirs du temps: Les visiteurs II  1998.0
```

### 2.2.3 Filtering text

To filter a pattern in text instead of the entire text, we can use: - LIKE - NOT LIKE - IN

Select the names of all people whose names begin with 'B'.

```
[91]: duckdb.query(\
        "select name from people where name like 'B%'" \
    ).to_df()
```

```
[91]:   name
0  B.J. Novak
1  Babak Najafi
2  Babar Ahmed
3  Bahare Seddiqi
4  Bai Ling
```

```

..      ...
440    Buster Keaton
441    Busy Philipps
442      Buzz Aldrin
443    Byron Howard
444      Byron Mann

[445 rows x 1 columns]

```

Select the names of people whose names have ‘r’ as the second letter.

```

[92]: duckdb.query(\
        "select name from people where name like '_r%'\n"
        ).to_df()

```

```

[92]:          name
0      Ara Celi
1    Aramis Knight
2  Arben Bajraktaraj
3  Arcelia Ram  rez
4    Archie Kao
..      ...
525    Troy Garity
526    Troy Miller
527    Troy Nixey
528  Ursula Andress
529    Wray Crawford

[530 rows x 1 columns]

```

Select the names of people whose names don’t start with ‘A’.

```

[93]: duckdb.query(\
        "select name from people where name not like 'A%'\n"
        ).to_df()

```

```

[93]:          name
0      50 Cent
1      lex Angulo
2    lex de la Iglesia
3      ngela Molina
4      B.J. Novak
..      ...
7763    Zohra Segal
7764  Zooey Deschanel
7765    Zoran Lisinac
7766    Zubaida Sahar
7767    Zuhair Haddad

```

[7768 rows x 1 columns]

Select the title and release\_year of all films released in 1990 or 2000 that were longer than two hours.

```
[94]: duckdb.query(\
        "select title, release_year from films where release_year in_
        ↪(1990, 2000) and duration > 120"\
        ).to_df()
```

```
[94]:
```

	title	release_year
0	Dances with Wolves	1990.0
1	Die Hard 2	1990.0
2	Ghost	1990.0
3	Goodfellas	1990.0
4	Mo' Better Blues	1990.0
5	Pretty Woman	1990.0
6	The Godfather: Part III	1990.0
7	The Hunt for Red October	1990.0
8	All the Pretty Horses	2000.0
9	Almost Famous	2000.0
10	Bamboozled	2000.0
11	Cast Away	2000.0
12	Chocolat	2000.0
13	Dancer in the Dark	2000.0
14	Erin Brockovich	2000.0
15	Finding Forrester	2000.0
16	Fiza	2000.0
17	Gladiator	2000.0
18	Gone in Sixty Seconds	2000.0
19	Keeping the Faith	2000.0
20	Love & Basketball	2000.0
21	Men of Honor	2000.0
22	Mission: Impossible II	2000.0
23	Pandaemonium	2000.0
24	Pay It Forward	2000.0
25	Pollock	2000.0
26	Proof of Life	2000.0
27	Quills	2000.0
28	Reindeer Games	2000.0
29	Space Cowboys	2000.0
30	The 6th Day	2000.0
31	The Contender	2000.0
32	The Family Man	2000.0
33	The House of Mirth	2000.0
34	The Legend of Bagger Vance	2000.0

35	The Patriot	2000.0
36	The Perfect Storm	2000.0
37	Thirteen Days	2000.0
38	Traffic	2000.0
39	Vertical Limit	2000.0
40	What Lies Beneath	2000.0
41	What Women Want	2000.0

Select the title and language of all films in English, Spanish, or French using IN.

```
[95]: duckdb.query(\
    "select title, language from films where language in ('English',
    ↪ 'Spanish', 'French')"\
    ).to_df()
```

```
[95]:
      title language
0    The Broadway Melody English
1      Hell's Angels English
2    A Farewell to Arms English
3      42nd Street English
4    She Done Him Wrong English
...
4742    The Blue Room  French
4743  Animal Kingdom: Let's go Ape French
4744    Evolution      French
4745  They Will Have to Kill Us First French
4746    Irreplaceable  French
```

[4747 rows x 2 columns]

Select the title, certification and language of all films certified NC-17 or R that are in English, Italian, or Greek.

```
[96]: duckdb.query(\
    "select title, certification, language from films where \
    certification in ('NC-17', 'R') and language in ('English',
    ↪ 'Italian', 'Greek')"\
    ).to_df()
```

```
[96]:
      title certification language
0    Pink Flamingos      NC-17 English
1    The Evil Dead      NC-17 English
2    Showgirls          NC-17 English
3    Orgazmo            NC-17 English
4    L.I.E.             NC-17 English
...
2001    The Neon Demon      R English
2002    The Perfect Match   R English
```

2003	The Purge: Election Year	R	English
2004	The Veil	R	English
2005	Triple 9	R	English

[2006 rows x 3 columns]

- Count the unique titles from the films database and use the alias provided (nineties\_english\_films\_for\_teens).
- Filter to include only movies with a release\_year from 1990 to 1999, inclusive.
- Add another filter narrowing your query down to English-language films.
- Add a final filter to select only films with 'G', 'PG', 'PG-13' certifications.

```
[97]: duckdb.query(\
        "select count(distinct(title)) as nineties_english_films_for_teens_
        ↪from films\
        where (release_year between 1990 and 1999)\
        and language = 'English'\
        and certification in ('G', 'PG', 'PG-13')"\
        ).to_df()
```

```
[97]:    nineties_english_films_for_teens
0                                           310
```

## 2.2.4 NULL Values

Select the title of every film that doesn't have a budget associated with it and use the alias no\_budget\_info.

```
[98]: duckdb.query(\
        "select title as no_budget_info from films where budget is null"\
        ).to_df()
```

```
[98]:    no_budget_info
0      Pandora's Box
1    The Prisoner of Zenda
2      The Blue Bird
3          Bambi
4      State Fair
..           ...
425    Unforgotten
426          Wings
427      Wolf Creek
428    Wuthering Heights
429 Yu-Gi-Oh! Duel Monsters
```

[430 rows x 1 columns]

Count the number of films with a language associated with them and use the alias

count\_language\_known.

```
[99]: duckdb.query(\
        "select count(*) as count_language_known from films where language_
        ↪is not null"\
        ).to_df()
```

```
[99]:      count_language_known
0              4955
```

## 2.3 Chapter 3: Aggregate Functions

Topics Covered:

- Summarizing Data
- Summarizing Subsets
- Aliasing and Arithmetic

### 2.3.1 Summarizing Data

- Count, MIN and MAX can work on various data types

Use the SUM() function to calculate the total duration of all films and alias with total\_duration.

```
[100]: duckdb.query(\
        "select sum(duration) as total_duration from films"\
        ).to_df()
```

```
[100]:      total_duration
0      534882.0
```

Calculate the average duration of all films and alias with average\_duration.

```
[101]: duckdb.query(\
        "select avg(duration) as average_duration from films"\
        ).to_df()
```

```
[101]:      average_duration
0      107.947931
```

Find the most recent release\_year in the films table, aliasing as latest\_year.

```
[102]: duckdb.query(\
        "select max(release_year) as latest_year from films"\
        ).to_df()
```

```
[102]:      latest_year
0      2016.0
```

Find the duration of the shortest film and use the alias `shortest_film`.

```
[103]: duckdb.query(\
        "select min(duration) as shortest_film from films"\
        ).to_df()
```

```
[103]:   shortest_film
0      7.0
```

### 2.3.2 Summarizing Subsets

Use `SUM()` to calculate the total gross for all films made in the year 2000 or later, and use the alias `total_gross`.

```
[104]: duckdb.query(\
        "select sum(gross) as total_gross from films\
        where release_year >= 2000"\
        ).to_df()
```

```
[104]:   total_gross
0  1.509009e+11
```

Calculate the average amount grossed by all films whose titles start with the letter 'A' and alias with `avg_gross_A`.

```
[105]: duckdb.query(\
        "select avg(gross) as avg_gross_A from films where title like 'A%'\
        ).to_df()
```

```
[105]:   avg_gross_A
0  4.789324e+07
```

Calculate the lowest gross film in 1994 and use the alias `lowest_gross`.

```
[106]: duckdb.query(\
        "select min(gross) as lowest_gross from films where release_year =\
        ↪1994"\
        ).to_df()
```

```
[106]:   lowest_gross
0    125169.0
```

Calculate the highest gross film between 2000 and 2012, inclusive, and use the alias `highest_gross`

```
[107]: duckdb.query(\
        "select max(gross) as highest_gross from films where release_year\
        ↪between 2000 and 2012"\
        ).to_df()
```



```
[107]: highest_gross
0      760505847.0
```

Calculate the average facebook\_likes to one decimal place and assign to the alias, avg\_facebook\_likes.

```
[108]: duckdb.query(\
        "select round(avg(facebook_likes), 1) as avg_facebook_likes from_
        ↪reviews"\
        ).to_df()
```

```
[108]: avg_facebook_likes
0      NaN
```

Calculate the average budget from the films table, aliased as avg\_budget\_thousands, and round to the nearest thousand.

```
[109]: duckdb.query(\
        "select round(avg(budget), -3) as avg_budget_thousands from films"\
        ).to_df()
```

```
[109]: avg_budget_thousands
0      39903000.0
```

### 2.3.3 Aliasing and arithmetic

- Arithmetic on same datatype gives the same datatype in SQL, eg. int divided by int gives int

Select the title and duration in hours for all films and alias as duration\_hours; since the current durations are in minutes, you'll need to divide duration by 60.0.

```
[110]: duckdb.query(\
        "select title, duration / 60.0 as duration_hours from films"\
        ).to_df()
```

```
[110]:
```

	title	duration_hours
0	Intolerance: Love's Struggle Throughout the Ages	2.050000
1	Over the Hill to the Poorhouse	1.833333
2	The Big Parade	2.516667
3	Metropolis	2.416667
4	Pandora's Box	1.833333
...	...	...
4963	Unforgotten	0.750000
4964	Wings	0.500000
4965	Wolf Creek	NaN
4966	Wuthering Heights	2.366667
4967	Yu-Gi-Oh! Duel Monsters	0.400000

[4968 rows x 2 columns]

Calculate the percentage of people who are no longer alive and alias the result as `percentage_dead`.

```
[111]: duckdb.query(\
        "SELECT count(deathdate) * 100.0 / count(id) AS percentage_dead_
        ↪FROM people;"\
        ).to_df()
```

```
[111]:    percentage_dead
0      9.372395
```

Find how many decades (period of ten years) the films table covers by using `MIN()` and `MAX()`; alias as `number_of_decades`.

```
[112]: duckdb.query(\
        "SELECT (MAX(release_year) - MIN(release_year)) / 10.0 AS_
        ↪number_of_decades FROM films;"\
        ).to_df()
```

```
[112]:    number_of_decades
0      10.0
```

## 2.4 Sorting and Grouping

### 2.4.1 Sorting Results

Select the name of each person in the people table, sorted alphabetically.

```
[113]: duckdb.query(\
        "SELECT name from people order by name"\
        ).to_df()
```

```
[113]:          name
0      50 Cent
1  A. Michael Baldwin
2    A. Raven Cruz
3    A.J. Buckley
4    A.J. DeLucia
...
8392  Ã"scar Jaenada
8393  Ã‰mile Gaudreault
8394  Ã‰milie Dequenne
8395  Ã‰ric Tessier
8396  Ã©tienne Faure

[8397 rows x 1 columns]
```

Select the title and duration for every film, from longest duration to shortest.

```
[114]: duckdb.query(\
        "SELECT title, duration from films order by duration desc"\
        ).to_df()
```

```
[114]:
```

	title	duration
0	Carlos	334.0
1	Blood In, Blood Out	330.0
2	Heaven's Gate	325.0
3	The Legend of Suriyothai	300.0
4	Das Boot	293.0
...	...	...
4963	Barfi	NaN
4964	Destiny	NaN
4965	Karachi se Lahore	NaN
4966	Romantic Schemer	NaN
4967	Wolf Creek	NaN

[4968 rows x 2 columns]

Select the release\_year, duration, and title of films ordered by their release year and duration, in that order.

```
[115]: duckdb.query(\
        "SELECT release_year, duration, title from films order by_\
        ↪release_year, duration"\
        ).to_df()
```

```
[115]:
```

	release_year	duration	title
0	1916.0	123.0	Intolerance: Love's Struggle Throughout the Ages
1	1920.0	110.0	Over the Hill to the Poorhouse
2	1925.0	151.0	The Big Parade
3	1927.0	145.0	Metropolis
4	1929.0	100.0	The Broadway Melody
...	...	...	...
4963	NaN	197.0	Deadline Gallipoli
4964	NaN	240.0	Emma
4965	NaN	286.0	The Company
4966	NaN	334.0	Carlos
4967	NaN	NaN	Wolf Creek

[4968 rows x 3 columns]

Select the certification, release\_year, and title from films ordered first by certification (alphabetically) and second by release year, starting with the most recent year.

```
[116]: duckdb.query(\
        "SELECT certification, release_year, title from films order by_\
        ↪certification, release_year desc"\
        ).to_df()
```

```
).to_df()
```

```
[116]:      certification  release_year      title
0      Approved      1967.0      You Only Live Twice
1      Approved      1967.0      Point Blank
2      Approved      1967.0      In Cold Blood
3      Approved      1966.0      Torn Curtain
4      Approved      1966.0  The Good, the Bad and the Ugly
...      ...      ...      ...
4963      NaN      NaN      Trapped
4964      NaN      NaN      Twisted
4965      NaN      NaN      Unforgettable
4966      NaN      NaN      Unforgotten
4967      NaN      NaN      Wings
```

```
[4968 rows x 3 columns]
```

## 2.4.2 Grouping Data

- To summarize / aggregate data for a specific group of results, we use group by

Select the `release_year` and count of films released in each year aliased as `film_count`.

```
[117]: duckdb.query(\
        "SELECT release_year, count(*) as film_count from films group by_\
        ↪release_year"\
        ).to_df()
```

```
[117]:      release_year  film_count
0      1916.0      1
1      1920.0      1
2      1925.0      1
3      1927.0      1
4      1929.0      2
..      ...      ...
87     2013.0     236
88     2014.0     252
89     2015.0     226
90     2016.0     106
91      NaN      42
```

```
[92 rows x 2 columns]
```

Select the `release_year` and average duration aliased as `avg_duration` of all films, grouped by `release_year`.

```
[118]: duckdb.query(\
```

```
"SELECT release_year, avg(duration) as avg_duration from films_
↳group by release_year"\
).to_df()
```

```
[118]:   release_year  avg_duration
0      1916.0      123.000000
1      1920.0      110.000000
2      1925.0      151.000000
3      1927.0      145.000000
4      1929.0      105.000000
..      ...      ...
87     2013.0      108.140426
88     2014.0      105.426295
89     2015.0      106.098214
90     2016.0      109.632075
91         NaN       77.439024
```

[92 rows x 2 columns]

Select the `release_year`, `country`, and the maximum budget aliased as `max_budget` for each year and each country; sort your results by `release_year` and `country`.

```
[119]: duckdb.query(\
        "SELECT release_year, country, max(budget) as max_budget from films_
↳group by release_year, country\
        order by release_year, country"\
).to_df()
```

```
[119]:   release_year  country  max_budget
0      1916.0      USA      385907.0
1      1920.0      USA      100000.0
2      1925.0      USA      245000.0
3      1927.0  Germany      6000000.0
4      1929.0  Germany         NaN
..      ...      ...      ...
500         NaN    Poland         NaN
501         NaN    Sweden         NaN
502         NaN      UK         NaN
503         NaN      USA      5000000.0
504         NaN     NaN         NaN
```

[505 rows x 3 columns]

### 2.4.3 Filtering grouped data (HAVING)

- Select `country` from the `films` table, and get the distinct count of certification aliased as `certification_count`.
- Group the results by `country`.

- Filter the unique count of certifications to only results greater than 10.

```
[120]: duckdb.query(\
        "SELECT country, count(distinct(certification)) as_
        ↪certification_count from films\
        group by country\
        having count(distinct(certification)) > 10"\
        ).to_df()
```

```
[120]:   country  certification_count
0      USA                      12
```

- Select the country and the average budget as average\_budget, rounded to two decimal, from films.
- Group the results by country.
- Filter the results to countries with an average budget of more than one billion (1000000000).
- Sort by descending order of the average\_budget.

```
[121]: duckdb.query(\
        "SELECT country, round(avg(budget), 2) as average_budget from films\
        group by country\
        having round(avg(budget), 2) > 1000000000\
        order by average_budget desc"\
        ).to_df()
```

```
[121]:   country  average_budget
0  South Korea  1.383960e+09
1    Hungary  1.260000e+09
```

### Bringing it all together - Final Exercise

- Select the release\_year for each film in the films table, filter for records released after 1990, and group by release\_year.

```
[122]: duckdb.query(\
        "SELECT release_year from films where release_year > 1990 group by_
        ↪release_year"\
        ).to_df()
```

```
[122]:   release_year
0      1991.0
1      1992.0
2      1993.0
3      1994.0
4      1995.0
5      1996.0
6      1997.0
```

7	1998.0
8	1999.0
9	2000.0
10	2001.0
11	2002.0
12	2003.0
13	2004.0
14	2005.0
15	2006.0
16	2007.0
17	2008.0
18	2009.0
19	2010.0
20	2011.0
21	2012.0
22	2013.0
23	2014.0
24	2015.0
25	2016.0

Modify the query to include the average budget aliased as `avg_budget` and average gross aliased as `avg_gross` for the results we have so far.

```
[123]: duckdb.query(\
        "SELECT release_year, avg(budget) as avg_budget, avg(gross) as \
        ↪avg_gross\
        from films\
        where release_year > 1990 group by release_year"\
    ).to_df()
```

```
[123]:
```

	release_year	avg_budget	avg_gross
0	1991.0	2.517655e+07	5.384450e+07
1	1992.0	2.598203e+07	6.366520e+07
2	1993.0	2.072979e+07	4.530209e+07
3	1994.0	2.901377e+07	5.939567e+07
4	1995.0	3.277500e+07	4.490952e+07
5	1996.0	3.162061e+07	4.204417e+07
6	1997.0	5.942449e+07	4.479377e+07
7	1998.0	4.046000e+07	3.837701e+07
8	1999.0	3.898178e+07	3.807218e+07
9	2000.0	3.493138e+07	4.217263e+07
10	2001.0	3.768731e+07	4.325572e+07
11	2002.0	3.259851e+07	4.351115e+07
12	2003.0	3.720865e+07	4.872775e+07
13	2004.0	4.686534e+07	4.072653e+07
14	2005.0	7.032394e+07	4.115914e+07
15	2006.0	9.396893e+07	3.923786e+07

16	2007.0	3.527113e+07	4.626750e+07
17	2008.0	4.180489e+07	4.457351e+07
18	2009.0	3.707329e+07	4.620744e+07
19	2010.0	4.609466e+07	4.990833e+07
20	2011.0	3.777525e+07	4.578584e+07
21	2012.0	4.133182e+07	6.287353e+07
22	2013.0	4.051904e+07	5.615836e+07
23	2014.0	3.532580e+07	6.241214e+07
24	2015.0	3.929833e+07	7.257330e+07
25	2016.0	5.664274e+07	7.692404e+07

Modify the query once more so that only years with an average budget of greater than 60 million are included.

```
[124]: duckdb.query(\
        "SELECT release_year, AVG(budget) AS avg_budget, AVG(gross) AS \
        ↪avg_gross\
        FROM films\
        WHERE release_year > 1990\
        GROUP BY release_year\
        having avg(budget) > 60000000"\
        ).to_df()
```

```
[124]:  release_year  avg_budget  avg_gross
0      2005.0  7.032394e+07  4.115914e+07
1      2006.0  9.396893e+07  3.923786e+07
```

Finally, order the results from the highest average gross and limit to one.

```
[125]: duckdb.query(\
        "SELECT release_year, AVG(budget) AS avg_budget, AVG(gross) AS \
        ↪avg_gross\
        FROM films\
        WHERE release_year > 1990\
        GROUP BY release_year\
        having avg(budget) > 60000000\
        order by avg_gross desc\
        limit 1"\
        ).to_df()
```

```
[125]:  release_year  avg_budget  avg_gross
0      2005.0  7.032394e+07  4.115914e+07
```

```
[ ]:
```